



Prompt Lab — Key Takeaways

Prompting for Leadership · Demystifying AI seminar

Synthetic scenarios only — no client or personal data

01 The six-part anatomy of a great prompt

A fantastic prompt covers all six. Use it as a checklist before you hit send.

■ Role

Who the AI should be — an analyst, an RM, an architect. Sets expertise and voice.

■ Context

The situation and audience — who's reading, what they already know.

■ Task

One clear, specific job to do.

■ Format

The exact structure you want back — sections, table, word limit.

■ Constraints / Tone

Length, register, and what to avoid.
The dial that matters most for people-facing text.

■ Examples

A short sample of the standard you expect.

02 What the exercises taught us

Rating the prompts

1. **Precision in, precision out.** A vague prompt earns a vague answer — the model mirrors the quality of your ask.
2. **More isn't better.** "Overcooked" prompts — piling on roles, contradictory limits ("20 tables in 150 words"), asking it to invent what's missing — backfire. Right-size the ask.
3. **Judge a prompt by its anatomy.** The winner isn't the longest or fanciest; it's the one that's specific, structured, and honest about uncertainty.

Writing the prompt

4. **Anchor it to the actual task.** Write the prompt *in context* to what's being asked or produced — the artefact, the reader, the decision it feeds.
5. **Specify the exact format** — numbered sections, a table, a word limit; the shape you want back.
6. **Tell it what **not** to do.** For BRDs, credit memos, legacy code: don't invent, don't resolve ambiguities, flag hidden rules with evidence, and say when data is stale rather than guess.

Tone is a business control, not a nicety

For **customer- and front-facing** communication the **tonality you specify in the prompt** drives how the output lands — warmth, formality, no jargon, no grovelling. Getting it wrong is a relationship risk, so state it explicitly. For **internal, engineering or AI/developer-led** prompts that same tonality matters far less — there, spend your words on **precision, evidence and structure** instead. In short:

customer-facing → set the tone

developer-led → optimise for precision

Overarching

- **Judgement stays with you.** AI drafts, you decide. Keep a human in the loop for risk, compliance and client-facing work.
- **Governance first.** Synthetic scenarios only — never enter client, personal or confidential data into consumer tools.

The reusable rule of thumb

- Say **who** it is, **who** it's for, the **one** job, the **shape** of the answer, the **limits**, and **show an example**.
- Then adjust **one** variable at a time — tone, or length, or format — rather than rewriting from scratch.

03 What you can change day-to-day — copilots & LLMs

Small habits that lift quality immediately, whether you're in ChatGPT, Claude, Microsoft 365 Copilot or Salesforce Einstein.

- **Open with role + context in one line.** *"You are a credit analyst writing a 3-minute pre-read for the risk committee..."*
- **State the format you want** before you ask — bullets, a table, five numbered sections, under 180 words.
- **Add two constraints:** one on **tone** (for people-facing text) and one on **honesty** — *"neutral tone; where data is stale, say so — do not guess."*
- **Give it the source and pin it there.** Paste the memo/notes/code and say *"use only the material provided; don't add anything."*
- **Iterate by one variable.** If the answer's off, change the format or the length or the tone — not everything at once.
- **Match effort to channel.** Customer-facing → set tone explicitly and review before sending.
Code/analysis → prioritise precision, evidence and citations over voice.
- **Ask for its assumptions and gaps.** *"List anything you were unsure about or had to assume"* surfaces risk you'd otherwise miss.
- **Never** paste client, personal or confidential data into a consumer tool — use the sanctioned enterprise copilot with the trust controls turned on.

04 Guardrails as configuration — LLMs embedded in an OEM

When the LLM lives inside Salesforce (Einstein / Agentforce) or Microsoft Dynamics 365 & Copilot, most guardrails are **platform configuration**, not code. Set them centrally so end users inherit them and can't strip them.

Ground the model to sanctioned data	Scope answers to approved objects/records only (Salesforce grounding & Data Cloud; Dynamics via Dataverse). Enforce row-/field-level security so the assistant inherits each user's permissions — least privilege, no elevated service accounts.
Own the instruction layer	Bake role, tone and do/don't rules into the platform's prompt templates (Einstein Prompt Builder; Copilot Studio topics & instructions) so they're applied server-side and users can't override them.
Turn on the trust / DLP layer	Salesforce Einstein Trust Layer: PII masking, toxicity detection, zero data-retention with the model vendor, full audit trail. Microsoft Purview DLP + sensitivity labels for Copilot; Azure OpenAI content filters set to the right severity.
Constrain the outputs	Require grounding/citations, enable toxicity & grounding scoring, and configure the assistant to defer — <i>"I don't have that information"</i> — instead of hallucinating.
Human-in-the-loop for actions	For anything customer-facing (case replies, emails) or that writes data, configure draft-only + approval workflow rather than auto-send. Scope connectors/agent actions to read-only where possible.

Tone & register by channel	Maintain separate prompt templates per channel — a formal, regulated tone for customer comms; a terser, precision-first template for internal/developer assistance. (This is the tonality lesson, encoded as config.)
Audit & retention	Enable prompt/response logging, set zero-retention with the LLM vendor, and apply your retention & monitoring policies. Review logs for drift and misuse.

Remember: the framework — Role · Context · Task · Format · Constraints/Tone · Examples — is the same whether you're typing into a chat box or configuring an enterprise copilot. In the copilot, you're just setting those defaults once, for everyone. · Prompt Lab · Prompting for Leadership · all scenarios synthetic.